

MAJOR RESEARCH QUESTIONS

What are the limits of crowdsourcing?

Research in crowdsourcing has spent the better part of a decade exploring ways of growing the limits of crowdsourcing and finding the boundaries of crowd work and microtasks; identifying challenges to this form of labor, overcoming them through novel designs of work-flows and processes, and repeating the process [4]. The question that has emerged among these researchers and through the work that they have produced then has been driving at *whether* there are limits to crowdsourcing (and, if so, what factors determine those limits). Through this lens, we can point to a number of contributions to the field that have extended the boundaries of crowd work.

The exploration of crowdsourcing's potential and limits has principally looked at manipulating and extending along three dimensions: 1) Decomposition (how high can we stack the salami?), 2) Work Abstraction, and 3) Volatile quality and availability? How much can you deal with?. We'll explore these aspects of crowdsourcing, discussing the extents to which work can be decomposed, contextually abstracted, and made more resilient to attrition of various forms. We'll also point to corresponding piecework literature addressing these aspects. Finally, we'll discuss how these elements will serve to constrain the upper and lower bounds of crowdsourcing as it relates to the question of the furthest limits of crowdsourcing.

Decomposition (how high can we stack the salami?)

Crowdsourcing has explored the decomposition of work for the better part of a decade; Kittur et al. specifically drive at this goal by addressing the possibility of crowdsourcing complex work [28]. Cheng et al. found that microtasks — though not necessarily *faster* than “macrotasks” — yield higher quality work, particularly when that work is susceptible to frequent interruptions [9].

Much of the research in the space of designing crowd work has sought to illustrate the potential to take highly creative or skilled work and generate high-quality results. Perhaps the most notable case study here can be found in Retelny et al.'s *Foundry*, which employed “flash teams” to accomplish work typically thought of as only within reach of experts [41].

Research exploring the decomposition of work more generally — that is, without the constraint of an employer/requester directing workers — has since emerged as well. Teevan, Liebling, and Lasecki further push the boundaries of decomposed work, exploring “selfsourcing”, and further this work with Teevan et al. [54, 55]. While some of this work doesn't strictly fall under “crowdsourcing”, the [scientific] management of the self as a worker — of sorts — will prove relevant as we trace the literature surrounding piecework.

The work described thus far has explored along myriad paths the ways that we can manage workers for the purpose of accomplishing complex tasks with rapidly sourced workers from all over the world. Critically, the work's defining characteristic involves the use of the Internet as a medium both to coordinate workers, as well as to do the work itself. But research into the “decomposition” of work more generally illustrates the same concepts of work that “Taylorism” and scientific management

sought to embody — Silberman, Irani, and Ross in particular foresaw this danger and warned of it in, but they were by no means alone in their warnings [29, 50, 37].

Following this thread of work, it may be said that much of the work in crowdsourcing has found ways of “vertically slicing” work such that each person is responsible for as small a task as can be made reasonable — taking a minute of audio, for instance, and transcribing that. The findings we take away from the crowdsourcing literature might be summed up as a number of ways to break tasks down [see 54, 41, 28] and measuring the costs as they relate to work completion & to a lesser extent to workers [see 8, 9, 29, 50].

Research into piecework has considered the decomposition of work; Graves describes piecework in 1981 as “... based on examination of various shop jobs, which included calculation of the standard time and compensation for each task. Most piecework plans required additional clerical and supervisory people (such as piecework clerks, inspectors, and “experts”)...”, allowing for what the *American Machinist* would later call “paper efficiency”, describing savings on paper that are in fact passed on to workers [18]. In this context, the introspection of tasks by experts and the careful work assignment & management of workers we see in historical piecework seem to parallel crowdsourcing.

identifying in analysis of data from the U.S. Bureau of Labor and Statistics that “larger establishment are less likely to use piece rates and predominantly female establishments more likely to do so; unionized establishments make significantly greater use of standard-rate pay than do nonunion establishments; **incentive pay is less likely in jobs with a variety of duties than in jobs with a narrow set of routinized duties**; and the ease of monitoring work quality is correlated with the use of incentive pay”

What's different today? We can turn to Graves's reflection on railroad piecework for a hint of what might have enabled the future of crowdwork in which we now live: “... only the largest and most wealthy railroads stood to gain substantially from scientific management, because only they had the resources to eliminate union resistance and *pay the overhead involved in installing work reorganization*” (emphasis added) [18].

We argue that crowdsourcing is more widespread today, and has the potential to become even larger in scale, because the underlying architecture of crowdwork is digital; reformulating workflows to decompose tasks and institute piecework compensation programs is easier today than it was for railroad companies which had to move physical capital to do so.

[a]2: Graves later argues that a significant obstacle to the introduction of piecework in railroad shops was the resistance of workers (the other being the resistance of management (for different and varying reasons)). In the cases of online platforms like AMT and digitally mediated platforms like Uber, where in both cases workers rarely if ever interact face-to-face and opportunities for coordination and collective action are severely limited, we should expect to see the challenges

stymieing collective action that we see in Salehi et al. and elsewhere [18, 46].]

Work Abstraction

Decomposition allows requesters to assign parallelized tasks without undue concern for the broader context of the work. The earlier example of transcribed audio is admittedly a carefully curated one — used to illustrate the ability to parallelize work across as many workers as there are tasks — but even in this case problems arise; breaking an hour-long interview into 3,600 1-second-long tasks would seem to make it impossible to do the tasks for lack of *context*, but the other extreme (one hour-long block) offers no performance improvement over conventionally sourced work.

The question, then, becomes about how to break tasks down as much as possible (as discussed in the previous section) without losing the necessary context to do a task. Stated another way, how can we design work so that as little context as possible is provided, without leaving out context or information necessary to complete the task? As Verroios and Bernstein frame it, “the crowd must be able to act with global understanding when each contributor only has access to local views” [60].

While some research has attempted to increase workers’ awareness of the work they do in apparent efforts to yield higher quality work, others have turned this constraint of crowdwork into a feature — the abstraction from the work itself appears to allow requesters to engage workers in work with — for instance — sensitive, confidential or otherwise personally identifying information. We will discuss these and other avenues of research here.

Cascade demonstrated that it’s possible to break certain classes of tasks apart in such a way that they yield taxonomies of various subjects, a task generally thought to be sufficiently complex that only expert workers with top-down awareness of the task — in this case, with awareness of all the constituent colors — could complete the task [10]. Verroios and Bernstein further illustrate this potential by forming a task one might consider highly contextually dependent — summarizing the contents of a movie — in such a way that crowd workers could contribute small pieces of work without needing to know the content of the rest of the project [60].

Volatile quality and availability? How much can you deal with?

A number of researchers have identified worker attrition, variability of worker performance, and uncertainty about good versus bad-faith actors as open questions of crowdwork [15, 22]. [a12: We can and should discuss the distinction between presumably “bad faith” workers & workers who are merely responding in kind to bad requesters — and the broader questions surrounding the roles that requesters as well as workers should play — but let it suffice to say that requesters have been trying to understand and manage what appears to them as inconsistent work. Their ways of responding to that variance in work quality has largely involved making the work more flexible and resilient to work (although some work has gone into investigating the causes, rather than treating the symptoms)]

Earlier we discussed Cheng et al.’s work measuring the impact that interruption has on worker performance [9]. This

work illustrates a broader sentiment in both the study and practice of crowdwork, that microtasks should be designed resiliently against the variability of workers, fully exploiting the abstracted nature of each piece of work [23, 31, 59]. That is to say, micro-tasks should be designed such that a single worker’s poor performance, or a good worker’s sudden departure, does not significantly impact the agenda of the work as a whole. While Cheng et al. found costs with breaking tasks into smaller components in the form of higher cumulative time to complete (albeit much shorter real time to complete, owing to parallelization), Lasecki et al. found that at least *some* performance can be recouped by stringing similar tasks together [9, 31, respectively].

Krishna et al. take a different approach; by “embracing error” and forming models describing the latency of workers in classifying objects at rapid speeds, the authors offer orders-of-magnitude improvements in various binary classification tasks [30]. And rather than building tasks to *tolerate* worker drop-off and attrition, some researchers have designed work predicated on the expectation of it instead: Celis et al. describe ways of assigning tasks in such a way that crowd workers would never be given enough information to piece together sensitive information about any single topic [7].

The work thus far seems to attempt to maximize the quality of work among workers through various means: 1) Identifying “bad” workers (fraught with problems as this characterization is) [15], 2) Designing tasks with break points to facilitate the on-boarding and off-boarding that happens anyway [9], and 3) Expecting certain levels of attrition and incorrectness and using that variability to their advantage [30].

Flexibility has been explored through the lens of Fordism, perhaps best illustrated by Tolliday and Zeitlin’s treatment describing turnover rates rising above 300% in the decade leading to the introduction of the assembly line in 1913. Specifically, the utilization of “... ‘semi-special’ machine tools which could be adapted [and] ... added flexibility through seasonal layoffs for production workers and the use of piece rates ... rather than a day wage system” [57].

In the field of piecework, the research covering this topic has both explored a breadth of tasks that might be rendered doable by piecemeal workers *as well as* longitudinally documented the success of these approaches. Here, we

Here, Hu’s work, saying of assembly line work that “it is assumed that men are of equal ability and every man can do any of the n jobs”, parallels the approach that dominated early research into crowd work — namely, using non-expert crowds for complex work [21]. This mindset in Hu’s analysis, and indeed the study of factory and mass manufacturing labor through the 20th century, substantively owes its existence to scientific management and the rigorous decomposition of work into tasks, discussed earlier, and persists to this day as it colors researchers’ goals and objectives in the study and design of crowd work.

Piecework’s influence on the abstraction of work into tasks, described above, is more than just caused by the decomposition of work; work abstraction itself makes it possible for

workers to come and go flexibly, prompting work requesters to consider ways to design these now discrete tasks in ways that maximize flexibility, both by allowing (and even anticipating) some inconsistency in worker availability *and* allowing and anticipating some inconsistency in the quality of the work output itself. It's to this area that we now turn our attention.

Piecework has seen work along this dimension spanning decades; Thompson investigate some of the ways that construction can benefit from the principles of scientific management. Thompson's thesis asserts that task work is predicated on the accurate scientific management of work, including the "miscellaneous tasks". Thompson argues — as early as 1913 — that "... one may be challenged to find any class of work involving labor either indoors or out-of-doors where tasks cannot be fixed by proper time-study" [56].

Broken down in this way, work could grow to unprecedented scales, but the quality of the work would remain relatively variable [36]. Textile work being a salient example, it took time for workers to acquire sufficient skill to do every aspect of the work so that the garment would be accepted by the company soliciting that work [61].

A compelling solution emerged in the early 20th century to break tasks down into discrete, manageable routines that could be taught relatively easily, and whose work output could be evaluated in abstraction from the rest of the work [3]. In Ford's assembly line, this meant that workers were not responsible for building a whole car, but a single very narrowly defined action that needed to be done on every car [34]. By the mid-20th century, Schoenberger writes, "... the intensification of the labor process is argued to have hit mental, physical, and social limits." [47].

This approach, "Fordism" (and its better-known contemporary "Taylorism" of similar ethos), can be seen today in crowd work and on-demand labor through the application of micro-tasks. Teevan, Iqbal, and Veh highlight some of the advantages of breaking work into pieces, facilitating evaluation and parallelization [53]. By decomposing and recomposing tasks, and in particular by assigning similarly natured work to the same workers, workers could become "experts" in a small aspect of the work that they did, speeding their work dramatically [31]. Perhaps more important, however, was that the breaking down of work into tasks has made it more practical to evaluate work at each stage [43].

So how does this affect crowdwork?

The work we've seen so far

- worst case: assembling iPhones (extant)
- average case: railroad workers and assembly lines
- high (complexity) case:

References

- [1] George A Akerlof. "The market for" lemons": Quality uncertainty and the market mechanism". In: *The quarterly journal of economics* (1970), pp. 488–500.
- [2] Andy Baio. *The Faces of Mechanical Turk*. Nov. 2008. URL: http://waxy.org/2008/11/the_faces_of_mechanical_turk/.
- [3] Peter Baker. "Production restructuring in the textiles and clothing industries". In: *New Technology, Work and Employment* 8.1 (1993), pp. 43–55. ISSN: 1468–005X. DOI: [10.1111/j.1468-005X.1993.tb00033.x](https://doi.org/10.1111/j.1468-005X.1993.tb00033.x). URL: <http://dx.doi.org/10.1111/j.1468-005X.1993.tb00033.x>.
- [4] Michael S. Bernstein et al. "Soylent: A Word Processor with a Crowd Inside". In: *UIST '10* (2010), pp. 313–322. DOI: [10.1145/1866029.1866078](https://doi.org/10.1145/1866029.1866078). URL: <http://doi.acm.org/10.1145/1866029.1866078>.
- [5] Robin Brewer, Meredith Ringel Morris, and Anne Marie Piper. "'Why Would Anybody Do This?': Understanding Older Adults' Motivations and Challenges in Crowd Work". In: *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*. CHI '16. New York, NY, USA: ACM, 2016, pp. 2246–2257. ISBN: 978–1-4503–3362–7. DOI: [10.1145/2858036.2858198](https://doi.org/10.1145/2858036.2858198). URL: <http://doi.acm.org/10.1145/2858036.2858198>.
- [6] Chantal Brisson et al. "Effect of duration of employment in piecework on severe disability among female garment workers". In: *Scandinavian journal of work, environment & health* (1989), pp. 329–334.
- [7] L. Elisa Celis et al. "Assignment Techniques for Crowdsourcing Sensitive Tasks". In: *Proceedings of the 19th ACM Conference on Computer-Supported Cooperative Work & Social Computing*. CSCW '16. New York, NY, USA: ACM, 2016, pp. 836–847. ISBN: 978–1-4503–3592–8. DOI: [10.1145/2818048.2835202](https://doi.org/10.1145/2818048.2835202). URL: <http://doi.acm.org/10.1145/2818048.2835202>.
- [8] Justin Cheng, Jaime Teevan, and Michael S. Bernstein. "Measuring Crowdsourcing Effort with Error-Time Curves". In: *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*. CHI '15. New York, NY, USA: ACM, 2015, pp. 1365–1374. ISBN: 978–1-4503–3145–6. DOI: [10.1145/2702123.2702145](https://doi.org/10.1145/2702123.2702145). URL: <http://doi.acm.org/10.1145/2702123.2702145>.
- [9] Justin Cheng et al. "Break it down: A comparison of macro-and microtasks". In: *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*. ACM. 2015, pp. 4061–4064.
- [10] Lydia B Chilton et al. "Cascade: Crowdsourcing taxonomy creation". In: *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. ACM. 2013, pp. 1999–2008.
- [11] Ben Craig and John Pencavel. "The behavior of worker cooperatives: The plywood companies of the Pacific Northwest". In: *The American Economic Review* (1992), pp. 1083–1105.
- [12] Ellen Cushing. *Dawn of the Digital Sweatshop*. Aug. 2012. URL: <http://www.eastbayexpress.com/oakland/dawn-of-the-digital-sweatshop/Content?oid=3301022>.
- [13] Bernhard Ebbinghaus and Jelle Visser. "When institutions matter: Union growth and decline in Western Europe, 1950–1995". In: *European Sociological Review* 15.2 (1999), pp. 135–158.

- [14] Karèn Fort, Gilles Adda, and K Bretonnel Cohen. "Amazon mechanical turk: Gold mine or coal mine?" In: *Computational Linguistics* 37.2 (2011), pp. 413–420.
- [15] Ujwal Gadiraju et al. "Understanding Malicious Behavior in Crowdsourcing Platforms: The Case of Online Surveys". In: *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*. CHI '15. New York, NY, USA: ACM, 2015, pp. 1631–1640. ISBN: 978–1-4503–3145–6. DOI: [10.1145/2702123.2702443](https://doi.org/10.1145/2702123.2702443). URL: <http://doi.acm.org/10.1145/2702123.2702443>.
- [16] Henry Laurence Gantt. *Work, wages, and profits*. Engineering Magazine Co., 1913.
- [17] Alan Gevins and Michael E Smith. "Neurophysiological measures of cognitive workload during human–computer interaction". In: *Theoretical Issues in Ergonomics Science* 4.1–2 (2003), pp. 113–131.
- [18] Carl Graves. "Applying Scientific Management Principles to Railroad Repair Shops — the Santa Fe Experience, 1904–18". In: *Business and Economic History* 10 (1981), pp. 124–136. ISSN: 08946825. URL: <http://www.jstor.org/stable/23702539>.
- [19] Russell Hardin. *Collective action*. Resources for the Future, 1982.
- [20] John Joseph Horton and Lydia B. Chilton. "The Labor Economics of Paid Crowdsourcing". In: *Proceedings of the 11th ACM Conference on Electronic Commerce*. EC '10. New York, NY, USA: ACM, 2010, pp. 209–218. ISBN: 978–1-60558–822–3. DOI: [10.1145/1807342.1807376](https://doi.org/10.1145/1807342.1807376). URL: <http://doi.acm.org/10.1145/1807342.1807376>.
- [21] Te C Hu. "Parallel Sequencing and Assembly Line Problems". In: *Operations Research* 9.6 (1961), pp. 841–848. DOI: [10.1287/opre.9.6.841](https://doi.org/10.1287/opre.9.6.841). eprint: <http://dx.doi.org/10.1287/opre.9.6.841>. URL: <http://dx.doi.org/10.1287/opre.9.6.841>.
- [22] Panagiotis G. Ipeirotis, Foster Provost, and Jing Wang. "Quality Management on Amazon Mechanical Turk". In: *Proceedings of the ACM SIGKDD Workshop on Human Computation*. HCOMP '10. Washington DC: ACM, 2010, pp. 64–67. ISBN: 978-1-4503-0222-7. DOI: [10.1145/1837885.1837906](https://doi.org/10.1145/1837885.1837906). URL: <http://doi.acm.org/10.1145/1837885.1837906>.
- [23] Shamsi T. Iqbal and Brian P. Bailey. "Effects of Intelligent Notification Management on Users and Their Tasks". In: *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. CHI '08. New York, NY, USA: ACM, 2008, pp. 93–102. ISBN: 978–1-60558–011–1. DOI: [10.1145/1357054.1357070](https://doi.org/10.1145/1357054.1357070). URL: <http://doi.acm.org/10.1145/1357054.1357070>.
- [24] Lilly Irani. "The cultural work of microwork". In: *New Media & Society* 17.5 (2015), pp. 720–739.
- [25] Lilly C. Irani and M. Six Silberman. "Stories We Tell About Labor: Turkopticon and the Trouble with "Design"". In: *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*. CHI '16. New York, NY, USA: ACM, 2016, pp. 4573–4586. ISBN: 978–1-4503–3362–7. DOI: [10.1145/2858036.2858592](https://doi.org/10.1145/2858036.2858592). URL: <http://doi.acm.org/10.1145/2858036.2858592>.
- [26] Lilly C. Irani and M. Six Silberman. "Turkopticon: Interrupting Worker Invisibility in Amazon Mechanical Turk". In: *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. CHI '13. New York, NY, USA: ACM, 2013, pp. 611–620. ISBN: 978–1-4503–1899–0. DOI: [10.1145/2470654.2470742](https://doi.org/10.1145/2470654.2470742). URL: <http://doi.acm.org/10.1145/2470654.2470742>.
- [27] Nicolas Kaufmann, Thimo Schulze, and Daniel Veit. "More than fun and money. Worker Motivation in Crowdsourcing—A Study on Mechanical Turk." In: *AMCIS*. Vol. 11. 2011, pp. 1–11.
- [28] Aniket Kittur et al. "CrowdForge: Crowdsourcing Complex Work". In: *Proceedings of the 24th Annual ACM Symposium on User Interface Software and Technology*. UIST '11. New York, NY, USA: ACM, 2011, pp. 43–52. ISBN: 978–1-4503–0716–1. DOI: [10.1145/2047196.2047202](https://doi.org/10.1145/2047196.2047202). URL: <http://doi.acm.org/10.1145/2047196.2047202>.
- [29] Aniket Kittur et al. "The Future of Crowd Work". In: *Proceedings of the 2013 Conference on Computer Supported Cooperative Work*. CSCW '13. New York, NY, USA: ACM, 2013, pp. 1301–1318. ISBN: 978–1-4503–1331–5. DOI: [10.1145/2441776.2441923](https://doi.org/10.1145/2441776.2441923). URL: <http://doi.acm.org/10.1145/2441776.2441923>.
- [30] Ranjay A. Krishna et al. "Embracing Error to Enable Rapid Crowdsourcing". In: *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*. CHI '16. New York, NY, USA: ACM, 2016, pp. 3167–3179. ISBN: 978–1-4503–3362–7. DOI: [10.1145/2858036.2858115](https://doi.org/10.1145/2858036.2858115). URL: <http://doi.acm.org/10.1145/2858036.2858115>.
- [31] Walter S. Lasecki et al. "The Effects of Sequence and Delay on Crowd Work". In: *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*. CHI '15. New York, NY, USA: ACM, 2015, pp. 1375–1378. ISBN: 978–1-4503–3145–6. DOI: [10.1145/2702123.2702594](https://doi.org/10.1145/2702123.2702594). URL: <http://doi.acm.org/10.1145/2702123.2702594>.
- [32] Min Kyung Lee et al. "Working with Machines: The Impact of Algorithmic and Data-Driven Management on Human Workers". In: *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*. CHI '15. New York, NY, USA: ACM, 2015, pp. 1603–1612. ISBN: 978–1-4503–3145–6. DOI: [10.1145/2702123.2702548](https://doi.org/10.1145/2702123.2702548). URL: <http://doi.acm.org/10.1145/2702123.2702548>.
- [33] Margaret Levi et al. "Union democracy reexamined". In: *Politics & Society* 37.2 (2009), pp. 203–228.
- [34] Alain Lipietz. "Towards Global Fordism?" In: *New Left Review* 0.132 (Mar. 1982). Last updated — 2013–02–24, p. 33. URL: <http://search.proquest.com/docview/1301937328?accountid=14026>.

- [35] Brian McInnis et al. "Taking a HIT: Designing Around Rejection, Mistrust, Risk, and Workers' Experiences in Amazon Mechanical Turk". In: *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*. CHI '16. New York, NY, USA: ACM, 2016, pp. 2271–2282. ISBN: 978–1–4503–3362–7. DOI: [10.1145/2858036.2858539](https://doi.org/10.1145/2858036.2858539). URL: <http://doi.acm.org/10.1145/2858036.2858539>.
- [36] Fergus Murray. "The decentralisation of production—the decline of the mass-collective worker?" In: *Capital & Class* 7.1 (1983), pp. 74–99.
- [37] Jeffrey V Nickerson. "Crowd work and collective learning". In: *Technology-Enhanced Professional Learning: Routledge, Forthcoming* (2013).
- [38] Oded Nov. "What Motivates Wikipedians?" In: *Commun. ACM* 50.11 (Nov. 2007), pp. 60–64. ISSN: 0001–0782. DOI: [10.1145/1297797.1297798](https://doi.org/10.1145/1297797.1297798). URL: <http://doi.acm.org/10.1145/1297797.1297798>.
- [39] Gabriele Paolacci, Jesse Chandler, and Panagiotis G Ipeirotis. "Running experiments on amazon mechanical turk". In: *Judgment and Decision making* 5.5 (2010), pp. 411–419.
- [40] Sidney Pollard. "Factory Discipline in the Industrial Revolution. 1". In: *The Economic History Review* 16.2 (1963), pp. 254–271.
- [41] Daniela Retelny et al. "Expert Crowdsourcing with Flash Teams". In: *Proceedings of the 27th Annual ACM Symposium on User Interface Software and Technology*. UIST '14. New York, NY, USA: ACM, 2014, pp. 75–85. ISBN: 978–1–4503–3069–5. DOI: [10.1145/2642918.2647409](https://doi.org/10.1145/2642918.2647409). URL: <http://doi.acm.org/10.1145/2642918.2647409>.
- [42] Jacob August Riis. *How the other half lives: Studies among the tenements of New York*. Penguin, 1901.
- [43] Jakob Rogstadius et al. "An Assessment of Intrinsic and Extrinsic Motivation on Task Performance in Crowdsourcing Markets." In: *ICWSM* 11 (2011), pp. 17–21.
- [44] Joel Ross et al. "Who Are the Crowdworkers?: Shifting Demographics in Mechanical Turk". In: *CHI '10 Extended Abstracts on Human Factors in Computing Systems*. CHI EA '10. New York, NY, USA: ACM, 2010, pp. 2863–2872. ISBN: 978–1–60558–930–5. DOI: [10.1145/1753846.1753873](https://doi.org/10.1145/1753846.1753873). URL: <http://doi.acm.org/10.1145/1753846.1753873>.
- [45] Donald F Roy. "Work satisfaction and social reward in quota achievement: An analysis of piecework incentive". In: *American Sociological Review* 18.5 (1953), pp. 507–514.
- [46] Niloufar Salehi et al. "We Are Dynamo: Overcoming Stalling and Friction in Collective Action for Crowd Workers". In: *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*. CHI '15. New York, NY, USA: ACM, 2015, pp. 1621–1630. ISBN: 978–1–4503–3145–6. DOI: [10.1145/2702123.2702508](https://doi.org/10.1145/2702123.2702508). URL: <http://doi.acm.org/10.1145/2702123.2702508>.
- [47] Erica Schoenberger. In: *Environment and Planning D: Society and Space* 6.3 (1988), pp. 245–262.
- [48] W Douglas Seymour. "Manual skills and industrial productivity". In: *Production Engineers Journal, Institution of* 33.4 (1954), pp. 240–248.
- [49] Aaron D. Shaw, John J. Horton, and Daniel L. Chen. "Designing Incentives for Inexpert Human Raters". In: *Proceedings of the ACM 2011 Conference on Computer Supported Cooperative Work*. CSCW '11. New York, NY, USA: ACM, 2011, pp. 275–284. ISBN: 978–1–4503–0556–3. DOI: [10.1145/1958824.1958865](https://doi.org/10.1145/1958824.1958865). URL: <http://doi.acm.org/10.1145/1958824.1958865>.
- [50] M. Six Silberman, Lilly Irani, and Joel Ross. "Ethics and Tactics of Professional Crowdsourcing". In: *XRDS* 17.2 (Dec. 2010), pp. 39–43. ISSN: 1528–4972. DOI: [10.1145/1869086.1869100](https://doi.org/10.1145/1869086.1869100). URL: <http://doi.acm.org/10.1145/1869086.1869100>.
- [51] Six Silberman. *Stop citing Ross et al. 2010, "Who are the crowdworkers?"*. Mar. 2015. URL: <https://medium.com/@silberman/stop-citing-ross-et-al-2010-who-are-the-crowdworkers-b3b9b1e8d300>.
- [52] Yongqiang Sun, Nan Wang, and Zeyu Peng. "Working for one penny: Understanding why people would like to participate in online tasks with low payment". In: *Computers in Human Behavior* 27.2 (2011). Web 2.0 in Travel and Tourism: Empowering and Changing the Role of Travelers, pp. 1033–1041. ISSN: 0747–5632. DOI: [http://dx.doi.org/10.1016/j.chb.2010.12.007](https://dx.doi.org/10.1016/j.chb.2010.12.007). URL: <http://www.sciencedirect.com/science/article/pii/S0747563210003742>.
- [53] Jaime Teevan, Shamsi T. Iqbal, and Curtis von Veh. "Supporting Collaborative Writing with Microtasks". In: *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*. CHI '16. New York, NY, USA: ACM, 2016, pp. 2657–2668. ISBN: 978–1–4503–3362–7. DOI: [10.1145/2858036.2858108](https://doi.org/10.1145/2858036.2858108). URL: <http://doi.acm.org/10.1145/2858036.2858108>.
- [54] Jaime Teevan, Daniel J. Liebling, and Walter S. Lasecki. "Selfsourcing Personal Tasks". In: *CHI '14 Extended Abstracts on Human Factors in Computing Systems*. CHI EA '14. New York, NY, USA: ACM, 2014, pp. 2527–2532. ISBN: 978–1–4503–2474–8. DOI: [10.1145/2559206.2581181](https://doi.org/10.1145/2559206.2581181). URL: <http://doi.acm.org/10.1145/2559206.2581181>.
- [55] Jaime Teevan et al. "Productivity Decomposed: Getting Big Things Done with Little Microtasks". In: *Proceedings of the 2016 CHI Conference Extended Abstracts on Human Factors in Computing Systems*. CHI EA '16. New York, NY, USA: ACM, 2016, pp. 3500–3507. ISBN: 978–1–4503–4082–3. DOI: [10.1145/2851581.2856480](https://doi.org/10.1145/2851581.2856480). URL: <http://doi.acm.org/10.1145/2851581.2856480>.
- [56] Sanford E Thompson. "Time-Study and Task Work". In: *The Journal of Political Economy* (1913), pp. 377–387.
- [57] Steven Tolliday and Jonathan Zeitlin. *Between fordism and flexibility*. Oxford, 1986.

- [58] HA Turner. "Trade unions, differentials and the leveling of wages". In: *The Manchester School* 20.3 (1952), pp. 227–282.
- [59] Rajan Vaish et al. "Low Effort Crowdsourcing: Leveraging Peripheral Attention for Crowd Work". In: *Second AAAI Conference on Human Computation and Crowdsourcing*. 2014.
- [60] Vasilis Verroios and Michael S Bernstein. "Context trees: Crowdsourcing global understanding from local views". In: *Second AAAI Conference on Human Computation and Crowdsourcing*. 2014.
- [61] Nicole Vezina, Daniel Tierney, and Karen Messing. "When is light work heavy? Components of the physical workload of sewing machine operators working at piecework rates". In: *Applied Ergonomics* 23.4 (1992), pp. 268–276.
- [62] Roger D Waldinger et al. "Helots no more: A case study of the Justice for Janitors campaign in Los Angeles". In: *The Ralph and Goldy Lewis Center for Regional Policy Studies* (1996).

Points to make:

- Blossoming of piecework
 - high point: consultants?
 - most cases: auto workers/etc...
 - worst case: sweatshops (especially in developing nations)
- what led to these outcomes?
 - “consultant” work came out well because the work was complex; this made it difficult to turn into a mass market commodity. We see consultants ranging skill levels like oDesk (implementing modules) to Accenture (on-demand teams of consultants).
 - auto workers, working in settings where capital couldn’t be moved as easily found themselves in the same workspace as a direct — if multi-stepped — result to the benefits of putting people in the same place to consolidate resources. Moreover, workers had leverage over factory owners as a result of that consolidated capital; operating equipment required training.
 - sweatshop workers fared the worst, for reasons that may seem obvious with hindsight. Source materials for textiles are easier to ship than mechanical components such as engines, making it easier for factories to relocate to developing nations (where cost-of-living, and consequently wages, would be lower). As wages, Co-LAs, and QoL rise, workers begin asking for (and later demanding) higher wages and better conditions. But where the auto workers have leverage, textile workers find only a precarious economic balance now tipped by their collective action, spurring manufacturers to move to a new locale

What forms of work design and worker management are viable?

- researchers have looked at how to increase worker productivity (e.g. finding the maximal speed at which gig workers can be expected to work before making errors) [8].
- we’ve also seen people “embrace error” [30].
- still other research has looked into ways to sandbox workers from the context of their work
- but scholarship looking into the design and management of work and workers isn’t new; lots of research into getting pieceworkers to do work more quickly [48].
- Researchers have even asked the age old question of *what motivates* pieceworkers (echoing similar research on Wikipedia and Mechanical Turk) [45, 38, 27]

What will work and the place of work look like for workers?

The metaphorical mechanics of these dynamics are still at play; workers and managers continue to interact in adversarial manners, despite substantive work into aligning the motivations of workers and requesters

The existing body of research has shed light on on-demand labor from various perspectives, and revealed a number of topics that, through our framing, are clearly situated together. Those topics are, at a high level, as follows:

1. the **processes** involved in making work into tasks, or discretization;
2. the outcomes (and indeed the **fallout**) of that discretization, both on the work itself as well as the workers; and finally
3. the **relationships** between workers and requesters of the work — both *cooperative* and *adversarial* cases.

The Fallout of Crowd Work

Irani and Silberman point out the disillusion that companies such as Amazon foster on platforms for work like AMT (see also Salehi et al.’s work continuing in the spirit of this observation to generate collective action to improve worker conditions) [26, 46]. Lee et al. find similarly that workers on gig work platforms are frustrated by the systems on which they work, to say little of the policies which these systems enforce [32].

We discussed the benefits of flexibility (both in the sense of having arbitrary workers perform tasks and in the sense that we can design tasks to be more resilient to poor work) in the previous section. It’s from that point in the literature that we turn our attention to the perhaps unintended effects of crowd work and the affordances for transience that we build into this mode of work. We’ll address two major areas of work under this subject: 1) Low Pay; and 2) Variable quality work.

Low Pay

Horton and Chilton identified problems with crowd work wages relatively early on, attempting to address this imbalance from a behavioral economic perspective — that is, identifying and presenting a model that describes a worker’s “*reservation wage*” [20]. This work has largely informed much of the research into and practice of estimating crowd work compensation [49, 39].

But we turn to Irani and Silberman’s discussion of “*Turkopticon*”, a system they designed to interrogate worker invisibility and to promote better wages across several dimensions [26]. Of particular relevance here, Irani and Silberman call to attention that “Turkers” are ultimately vulnerable to wage theft and pay rates that translate to well under minimum wage. Returning to Horton and Chilton, we find that the median “reservation wage” in 2010 was \$1.38, while the mean was \$3.63 [20].

Understanding workers’ motivations given these conditions has thus become a goal for some researchers [5]. Sun, Wang, and Peng conclude that “... solvers participate in online tasks not only for money but also for enjoyment or the sense of self-worth” [52]. This might have rung true in 2011, and certainly corroborates Ross et al.’s findings after investigating “who are the crowdworkers”, but as Silberman points out “we [have since] learned that most tasks on AMT are done by a small group of professional Turkers...” [44, 51].

Now, Irani and Silberman and later Salehi et al. cite insufficient pay as a central point of frustration among workers, via Irani and Cushing’s contributions in this space [46, 24, 12, 26].

On-demand workers were not the first to be exploited along the dimension of low pay rates. Frustration over low (and declining) pay was one of the chief grievances among then nascent British labor unions in the early 20th century [58]. This, Ebbinghaus and Visser argued, fueled the rocketing

union membership rates through the mid-20th century until 1980 (to which we'll return when we discuss Levi et al.'s reexamination of labor unions) [13, 33]. This realization has similarly fueled a body of research into the various incentive structures available to piecework employers [45].

The parallels between the complaints of low pay among crowd workers and other on-demand workers and the pieceworkers and later factory workers in the 20th century are inescapable. We argue further that the *causes* here — work decomposition, work abstraction, and flexibility — lead inexorably to low and declining pay for workers. Moreover, we point out that low pay leads to other negative outcomes both in on-demand work as well as in piecework and on assembly lines.

Variable quality work

Researchers have struggled with what we might generously call work of “variable quality” along two dimensions. The first, to use the characterization of one of these contributions, we can call “understanding malicious behavior” [15]. While some work has cast workers as “malicious” or at least adversarial parties, the evidence thus far suggests that workers behave in unexpected ways as they attempt to assert some control over their interaction with the system (a topic of discussion to which we'll return later) [32]. The second dimension of research in this space generally attempts to eke out the highest quality work possible from workers given the apparent difficulty in predicting work outcomes [30].

The effect low wages have had on piece work and factory workers is well-known; Gantt discuss this exact mechanism in his book on “... where there is no union, the class wage is practically gauged by the wages the poor workman will accept, and the good workman soon becomes discouraged and *sets his pace by that of his less efficient neighbor*, with the result that the general tone of the shop is lowered” (emphasis added) [16].

This research is similar to, but subtly different from, the notion of the “market for ‘lemons’” which Fort, Adda, and Cohen discuss; specifically, Akerlof's writing of a “market for ‘lemons’” describes a marketplace where the quality of the product or service is unknown to the buyer [14, 1]. The effect of this *perceived* uncertainty is that the *actual* trustworthiness drops precipitously as all of the consistent, reliable, high-quality workers capable of leaving these markets do so, leaving only the ones who cannot or will not establish their trustworthiness.

Relationships Between Workers and Managers

Suffice it to say that poor pay and poor work are linked, and that we should not be surprised to find this relationship play out online as strongly as it does offline. But the poor treatment of workers by managers — both human and algorithmic — do more than affect the economic relationships between workers and employers. Here, then, we turn to examine this facet of on-demand work and how these dynamics strikingly replicate the relationships researchers in labor advocacy encountered in the study of piecework and factory work.

This topic can be condensed into two major areas: 1) external (scientific) management, and the evaluation of workers as functional modules; and 2) the consequential resistance

workers express due to their perceived alienation and distance from managing forces.

External Management

We discussed Fordism and Taylorism earlier in our discussions of Decomposition (how high can we stack the salami?) and Volatile quality and availability? How much can you deal with?, but here the core of these paradigmatic views — the scientific management of work — becomes relevant. We use “external” here instead of “scientific”, however, to more broadly capture the disconnect between managers and workers. By describing it as thus, we can touch on the relationship that workers have with *researchers*, as well, even though that work is not strictly — or just not exclusively — of the same nature as the management and experience as when interacting with requesters.

First, intuitively, the variable-quality work we discussed previously has led to a large and growing body of research attempting to evaluate workers' performance and error rates across numerous dimensions; for example, Cheng, Teevan, and Bernstein explore the error rates of workers by operating on a sliding scale giving workers varying amounts of time to accomplish micro-tasks [8]. Irani and Silberman describe the treatment of workers as sorts of “human APIs” that can, importantly, be rigorously evaluated [25]. Gevins and Smith began to explore the neurophysiological effects of cognitively demanding tasks on workers, informing crowdsourcing research by suggesting the use of cognitive load assessments such as NASA Task Load Index surveys to evaluate workers pre and post-tasks [30, 8].

External management comes in other forms than scientific, as previously mentioned. Researchers in particular have noticed that their relationships with on-demand workers are, at the least, complex. Irani and Silberman point out that their relationships with Turkers are highly complex; specifically, their interactions with field sites in which they work as designers and mediators of change influence the relationships they have with Turkers [25].

The scientific management of pieceworkers has been well-studied under the umbrella of assembly line research, and even physiological study of pieceworkers closely resembles the research into cognitive loads and stress levels that we discussed among on-demand crowd workers [21, 6]. Even the complicated relationships between observers and workers themselves are not necessarily new; Riis's photodocumentary of pieceworkers has even been re-examined through an exercise asking crowd workers to photograph themselves for similar purposes as Riis's — to document and humanize an otherwise abstracted, invisible workforce [2, 26, 42].

Similarly, Pollard's words on the punishment factory workers faced — for example, that “unsatisfactory work was punished ... by fines or by dismissal” — seems especially relevant given the fears we now know to be ubiquitous on platforms such as AMT, Uber, and other on-demand markets [40, 32, 46, 26, 35].

Resistance

It shouldn't surprise us, then, that workers have resisted the management imposed on them both by other people and their systems, often without recourse or opportunity for feedback, let alone substantive input. Indeed, Lee et al. discover of Uber drivers that many toggle their availability to avoid being dispatched to more distant locations, resisting the intent of the designers of the systems and their "algorithmic and data-driven management" [32].

Resistance has sometimes been more coordinated, as well; we see this in Irani and Silberman's coverage on *Turkopticon* as workers collectively accumulated information about requesters, and in Salehi et al.'s work on *Dynamo*, which generated "Guidelines for Academic Requesters" written by crowd workers [26, 46].

Resistance against managers in piecework and factory labor settings are deeply well-explored, but perhaps the most relevant case study to draw on here is to be found in Waldinger et al.'s case study of "Justice for Janitors", where marginalized workers managed to raise awareness for their plight and secure support for badly needed reforms [62]. The achievements of labor advocacy groups such as labor unions as resistant, even adversarial organizations counter-balancing the management is somewhat well-understood [19, 11]. We argue that these threads of resistance against management in various forms are in fact one.