

Decomposing Work

Crowd work's perspective. **The crowdsourcing research into work decomposition has largely focused on minimizing the additional context necessary to do tasks, and making it easier to do tasks with less time.** This first thread is perhaps best described by Verroios and Bernstein as making crowd workers "... able to act with global understanding when each contributor only has access to local views" [24]. With the exception of a few cases (specifically, Kinnaird, Dabbish, and Kiesler's work which finds that greater work context fosters more reliably high-quality work), the micro task paradigm has emerged as the overwhelming favorite [20, 21, 6, 15].

As the additional context necessary to complete a task diminishes, the marginal cost of finding and doing tasks has increasingly become the focus of research. Chilton et al. illustrate the challenges on AMT, and some work has gone into ameliorating the problems specific to this work site (*Re-Launcher*), while other work designs tasks around gap time (*TwitCh Crowdsourcing & Wait-Learning*) [7, 17, 23, 4]. Cosley et al. attempts to address this by directing workers to tasks through "intelligent task routing" [8]. Much of this work and the work at the periphery of this space, then, has focused on minimizing the amount of time that people need to spend doing anything other than the work for which they are paid.

Earlier we discussed Cheng et al.'s work measuring the impact that interruption has on worker performance [6]. This work illustrates a broader sentiment in both the study and practice of crowd work, that microtasks should be designed resiliently against the variability of workers, fully exploiting the abstracted nature of each piece of work [12, 18, 22]. That is to say, micro-tasks should be designed such that a single worker's poor performance, or a good worker's sudden departure, does not significantly impact the agenda of the work as a whole. While Cheng et al. found costs with breaking tasks into smaller components in the form of higher cumulative time to complete (albeit much shorter real time to complete, owing to parallelization), Lasecki et al. found that at least *some* performance can be recouped by stringing similar tasks together [6, 18, respectively].

Yet more work looks at the general framing of tasks, chaining and arranging them to maximally exploit the attention and stress threshold of workers [3]. Rather than attempt to minimize the error rates in micro-tasks, as Kinnaird, Dabbish, and Kiesler suggested, we as a community have leaned *into* the peril of low-context work, "embracing error" in crowdsourcing [16].

Piecework's perspective. **The research community relating to piecework and labor has been wrestling with the decomposition of work for centuries.** The beginnings of systematic task decomposition stretch back as far as the 19th century, when Airy employed young boys at the Greenwich Observatory who "possessed the basic skills of mathematics, including 'Arithmetic, the use of Logarithms, and Elementary Algebra'" to compute, by hand, astronomical phenomena [10]. These workers became the first *computers*.

The work Airy solicited was interesting for several reasons. First, work output was quickly verifiable; Airy could assign variably skilled workers to compute values, and have other workers check their work. Second, tasks were discrete — that is, independent from one another. Finally, knowledge of the full scope of the project — indeed, knowledge of anything more than the problem set at hand — was unnecessary.

The insight of breaking tasks down into smaller components didn't find its audience until the early 20th century, with the rise of Fordism and scientific management (or Taylorism). From scientific management, we found that we could measure work at unprecedented resolution and precision. As Brown points out, piecework most greatly benefits the instrumented measurement of workers, but certainly in Ford and Taylor's time — and certainly in Airy's time — highly instrumented, automatic measurement of workers was all but impossible. As a result, the distillation of work into smaller chunks ultimately reached a limit of usefulness.

Piecework researchers address a number of the problems with the decomposition of work, and the conflicting pressures managers and workers put forth. Bewley in particular points out that the approach of paying workers by the piece is "... not practical for workers doing many tasks, because of the cost of establishing the rates and because piecework does not compensate workers for time spent switching tasks". Ultimately, Bewley argues that "[piecework is] infeasible, because ... total output is the joint product of varying groups of people" [1]. [a12: This is becoming both a complexity argument and a decomposition argument. ... I'm okay with using Bewley in both, but I want people to read this paper and not be like "wow this author sure got their money's worth from those three books!"]

What's different about crowd work. A number of factors in crowd work are different from piecework, chief among them being the relative ease with which the metaphorical "assembly line" can be changed. Computers make it possible to switch from one task to another unlike any arbitrary manufacturing factory possibly could; a worker could do any number of different *types* of tasks in the span of just a few minutes, driven in particular by the power Lessig points to — that system-designers can share, modify, and instantiate environments like sites of labor in a few lines of code [18, 19]. This has spurred an entire body of work investigating the effects of ordering, pacing, interruptions, and other factors in piecework that would have been all but impossible to measure consistently as few as 20 years ago [6, 5, 16].

Further, we've sliced work to such small scales that the marginal activities — things like finding work, cognitive task switching, etc. — have become relatively large compared to the tasks themselves [7]. In the historical case of piecework, moving metallurgical tools, mining equipment, or other industry materials would have been prohibitively difficult and slow; workers were encouraged to specialize in a single set of tasks, allowing pieceworkers to sequence their tasks optimally on their own [11].

Rather than fall into the trap that Irani warns of, — one which where crowd workers are rendered as “modular, protocol-defined computational services” — we may yield better results from crowd work if we think of workers as similar to specialized, repurposable tools [13]. [al2: feeling meh about this argument...]

Finally, instrumentation has reached a sufficiently advanced and ubiquitous point that the dream of scientific management and Taylorism — to measure every motion at every point throughout the workday and beyond — is not only doable, but trivial [25]. One of the major challenges Graves cites as preventing scientific management from being fully utilized, the difficulty of tracking work & workers, no longer exists [9].

Implications for crowd work research. crowd work research today is on the right track to investigate pipelining and meta-task design. That is, investigating better work discovery methods, producing tools for workers to make more informed decisions [see, for example, 14]. It’s not clear how much benefit there is in the further decomposition of work, given that we’ve hit bottlenecks with the cognitive stresses of switching between tasks as Lasecki et al. highlight [18].

References

- [1] Truman F Bewley. *Why wages don’t fall during a recession*. Harvard University Press, 1999.
- [2] Charles Brown. “Firms’ Choice of Method of Pay”. In: *Industrial & Labor Relations Review* 43.3 (1990), 165S–182S. doi: [10.1177/001979399004300311](https://doi.org/10.1177/001979399004300311). eprint: <http://ilr.sagepub.com/content/43/3/165S.full.pdf+html>. URL: <http://ilr.sagepub.com/content/43/3/165S.abstract>.
- [3] Carrie J. Cai, Shamsi T. Iqbal, and Jaime Teevan. “Chain Reactions: The Impact of Order on Microtask Chains”. In: *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*. CHI ’16. New York, NY, USA: ACM, 2016, pp. 3143–3154. ISBN: 978-1-4503-3362-7. doi: [10.1145/2858036.2858237](https://doi.org/10.1145/2858036.2858237). URL: <http://doi.acm.org/10.1145/2858036.2858237>.
- [4] Carrie J. Cai et al. “Wait-Learning: Leveraging Wait Time for Second Language Education”. In: *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*. CHI ’15. Seoul, Republic of Korea: ACM, 2015, pp. 3701–3710. ISBN: 978-1-4503-3145-6. doi: [10.1145/2702123.2702267](https://doi.org/10.1145/2702123.2702267). URL: <http://doi.acm.org/10.1145/2702123.2702267>.
- [5] Justin Cheng, Jaime Teevan, and Michael S. Bernstein. “Measuring Crowdsourcing Effort with Error-Time Curves”. In: *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*. CHI ’15. New York, NY, USA: ACM, 2015, pp. 1365–1374. ISBN: 978-1-4503-3145-6. doi: [10.1145/2702123.2702145](https://doi.org/10.1145/2702123.2702145). URL: <http://doi.acm.org/10.1145/2702123.2702145>.
- [6] Justin Cheng et al. “Break it down: A comparison of macro-and microtasks”. In: *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*. ACM. 2015, pp. 4061–4064.
- [7] Lydia B. Chilton et al. “Task Search in a Human Computation Market”. In: *Proceedings of the ACM SIGKDD Workshop on Human Computation*. HCOMP ’10. New York, NY, USA: ACM, 2010, pp. 1–9. ISBN: 978-1-4503-0222-7. doi: [10.1145/1837885.1837889](https://doi.org/10.1145/1837885.1837889). URL: <http://doi.acm.org/10.1145/1837885.1837889>.
- [8] Dan Cosley et al. “SuggestBot: Using Intelligent Task Routing to Help People Find Work in Wikipedia”. In: *Proceedings of the 12th International Conference on Intelligent User Interfaces*. IUI ’07. Honolulu, Hawaii, USA: ACM, 2007, pp. 32–41. ISBN: 1-59593-481-2. doi: [10.1145/1216295.1216309](https://doi.org/10.1145/1216295.1216309). URL: <http://doi.acm.org/10.1145/1216295.1216309>.
- [9] Carl Graves. “Applying Scientific Management Principles to Railroad Repair Shops — the Santa Fe Experience, 1904–18”. In: *Business and Economic History* 10 (1981), pp. 124–136. ISSN: 08946825. URL: <http://www.jstor.org/stable/23702539>.
- [10] David Alan Grier. *When computers were human*. Princeton University Press, 2013.
- [11] Robert A Hart and J Elizabeth Roberts. “The rise and fall of piecework–timework wage differentials: market volatility, labor heterogeneity, and output pricing”. In: (2013).
- [12] Shamsi T. Iqbal and Brian P. Bailey. “Effects of Intelligent Notification Management on Users and Their Tasks”. In: *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. CHI ’08. New York, NY, USA: ACM, 2008, pp. 93–102. ISBN: 978-1-60558-011-1. doi: [10.1145/1357054.1357070](https://doi.org/10.1145/1357054.1357070). URL: <http://doi.acm.org/10.1145/1357054.1357070>.
- [13] Lilly Irani. “The cultural work of microwork”. In: *New Media & Society* 17.5 (2015), pp. 720–739.
- [14] Lilly C. Irani and M. Six Silberman. “Turkopticon: Interrupting Worker Invisibility in Amazon Mechanical Turk”. In: *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. CHI ’13. New York, NY, USA: ACM, 2013, pp. 611–620. ISBN: 978-1-4503-1899-0. doi: [10.1145/2470654.2470742](https://doi.org/10.1145/2470654.2470742). URL: <http://doi.acm.org/10.1145/2470654.2470742>.
- [15] Peter Kinnaird, Laura Dabbish, and Sara Kiesler. “Workflow Transparency in a Microtask Marketplace”. In: *Proceedings of the 17th ACM International Conference on Supporting Group Work*. GROUP ’12. Sanibel Island, Florida, USA: ACM, 2012, pp. 281–284. ISBN: 978-1-4503-1486-2. doi: [10.1145/2389176.2389219](https://doi.org/10.1145/2389176.2389219). URL: <http://doi.acm.org/10.1145/2389176.2389219>.
- [16] Ranjay A. Krishna et al. “Embracing Error to Enable Rapid Crowdsourcing”. In: *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*. CHI ’16. New York, NY, USA: ACM, 2016, pp. 3167–3179. ISBN: 978-1-4503-3362-7. doi: [10.1145/2858036.2858115](https://doi.org/10.1145/2858036.2858115). URL: <http://doi.acm.org/10.1145/2858036.2858115>.

- [17] Pavel Kucherbaev et al. “ReLauncher: Crowdsourcing Micro-Tasks Runtime Controller”. In: *Proceedings of the 19th ACM Conference on Computer-Supported Cooperative Work & Social Computing*. CSCW ’16. New York, NY, USA: ACM, 2016, pp. 1609–1614. ISBN: 978–1-4503–3592–8. DOI: [10.1145/2818048.2820005](https://doi.org/10.1145/2818048.2820005). URL: <http://doi.acm.org/10.1145/2818048.2820005>.
- [18] Walter S. Lasecki et al. “The Effects of Sequence and Delay on Crowd Work”. In: *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*. CHI ’15. New York, NY, USA: ACM, 2015, pp. 1375–1378. ISBN: 978–1-4503–3145–6. DOI: [10.1145/2702123.2702594](https://doi.org/10.1145/2702123.2702594). URL: <http://doi.acm.org/10.1145/2702123.2702594>.
- [19] Lawrence Lessig. *Code*. Lawrence Lessig, 2006.
- [20] Jaime Teevan, Daniel J. Liebling, and Walter S. Lasecki. “Selfsourcing Personal Tasks”. In: *CHI ’14 Extended Abstracts on Human Factors in Computing Systems*. CHI EA ’14. New York, NY, USA: ACM, 2014, pp. 2527–2532. ISBN: 978–1-4503–2474–8. DOI: [10.1145/2559206.2581181](https://doi.org/10.1145/2559206.2581181). URL: <http://doi.acm.org/10.1145/2559206.2581181>.
- [21] Jaime Teevan et al. “Productivity Decomposed: Getting Big Things Done with Little Microtasks”. In: *Proceedings of the 2016 CHI Conference Extended Abstracts on Human Factors in Computing Systems*. CHI EA ’16. New York, NY, USA: ACM, 2016, pp. 3500–3507. ISBN: 978–1-4503–4082–3. DOI: [10.1145/2851581.2856480](https://doi.org/10.1145/2851581.2856480). URL: <http://doi.acm.org/10.1145/2851581.2856480>.
- [22] Rajan Vaish et al. “Low Effort Crowdsourcing: Leveraging Peripheral Attention for Crowd Work”. In: *Second AAAI Conference on Human Computation and Crowdsourcing*. 2014.
- [23] Rajan Vaish et al. “Twitch Crowdsourcing: Crowd Contributions in Short Bursts of Time”. In: *Proceedings of the 32nd Annual ACM Conference on Human Factors in Computing Systems*. CHI ’14. Toronto, Ontario, Canada: ACM, 2014, pp. 3645–3654. ISBN: 978–1-4503–2473–1. DOI: [10.1145/2556288.2556996](https://doi.org/10.1145/2556288.2556996). URL: <http://doi.acm.org/10.1145/2556288.2556996>.
- [24] Vasilis Verroios and Michael S Bernstein. “Context trees: Crowdsourcing global understanding from local views”. In: *Second AAAI Conference on Human Computation and Crowdsourcing*. 2014.
- [25] Emily Waltz. “How I quantified myself”. In: *Spectrum, IEEE* 49.9 (2012), pp. 42–47.